

# Transfer Learning

November 9, 2025

```
[1]: from torch.utils.data import Dataset, DataLoader
from torchvision.transforms import ToTensor, transforms, v2
from pathlib import Path
from PIL import Image
from torch import nn
import torch
import random
import matplotlib.pyplot as plt
import torchvision
from torchvision import datasets
from torchinfo import summary
```

```
[2]: device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
[3]: def accuracy_fn(y_pred, y_true):
    correct_num = (y_pred == y_true).sum()
    acc = correct_num / len(y_true) * 100
    return acc
```

```
[4]: def train_step(dataloader, model, cost_fn, optimizer, accuracy_fn, device):
    train_cost = 0
    train_acc = 0
    for batch, (x, y) in enumerate(dataloader):

        x = x.to(device)
        y = y.to(device)

        model.train()
        y_pred = model(x)
        cost = cost_fn(y_pred, y)
        train_acc += accuracy_fn(y_pred.argmax(dim=1), y)

        train_cost += cost

        optimizer.zero_grad()
        cost.backward()
        optimizer.step()
```

```

train_cost /= len(dataloader)
train_acc /= len(dataloader)

return train_cost, train_acc

def test_step(dataloader, model, cost_fn, accuracy_fn, device):
    test_cost = 0
    test_acc = 0
    model.eval()
    with torch.inference_mode():
        for x, y in dataloader:
            x = x.to(device)
            y = y.to(device)
            test_pred = model(x)

            test_cost += cost_fn(test_pred,y)
            test_acc += accuracy_fn(test_pred.argmax(dim=1), y)

    test_cost /= len(dataloader)
    test_acc /= len(dataloader)

    return test_cost, test_acc

```

## 0.1 Model

[5]:

```

weights = torchvision.models.EfficientNet_B1_Weights.DEFAULT
model = torchvision.models.efficientnet_b1(weights = weights)
model.to(device)

```

[5]:

```

EfficientNet(
  (features): Sequential(
    (0): Conv2dNormActivation(
      (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
      bias=False)
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
      track_running_stats=True)
      (2): SiLU(inplace=True)
    )
    (1): Sequential(
      (0): MBConv(
        (block): Sequential(
          (0): Conv2dNormActivation(
            (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1,
            1), groups=32, bias=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
            track_running_stats=True)
            (2): SiLU(inplace=True)

```

```

)
(1): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(8, 32, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
(2): Conv2dNormActivation(
    (0): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(stochastic_depth): StochasticDepth(p=0.0, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=16, bias=False)
            (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(16, 4, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(4, 16, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (2): Conv2dNormActivation(
            (0): Conv2d(16, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
)
(stochastic_depth): StochasticDepth(p=0.008695652173913044, mode=row)
)
)
(2): Sequential(
    (0): MBConv(
        (block): Sequential(
            (0): Conv2dNormActivation(
                (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (2): SiLU(inplace=True)
)
(1): Conv2dNormActivation(
    (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=96, bias=False)
        (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
)
(2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(96, 4, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(4, 96, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
(3): Conv2dNormActivation(
    (0): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(stochastic_depth): StochasticDepth(p=0.017391304347826087, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=144, bias=False)
            (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
)

```

```

        (3): Conv2dNormActivation(
            (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.026086956521739136, mode=row)
)
(2): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=144, bias=False)
            (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.034782608695652174, mode=row)
)
)
(3): Sequential(
    (0): MBConv(
        (block): Sequential(
            (0): Conv2dNormActivation(
                (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)

```

```

)
(1): Conv2dNormActivation(
    (0): Conv2d(144, 144, kernel_size=(5, 5), stride=(2, 2), padding=(2,
2), groups=144, bias=False)
        (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
)
(2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(144, 6, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(6, 144, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
(3): Conv2dNormActivation(
    (0): Conv2d(144, 40, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
(stochastic_depth): StochasticDepth(p=0.043478260869565216, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2), groups=240, bias=False)
            (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(stochastic_depth): StochasticDepth(p=0.05217391304347827, mode=row)
)
(2): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(240, 240, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2), groups=240, bias=False)
            (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(240, 40, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(40, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.06086956521739131, mode=row)
)
)
(4): Sequential(
    (0): MBConv(
        (block): Sequential(
            (0): Conv2dNormActivation(
                (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
            (1): Conv2dNormActivation(

```

```

        (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=240, bias=False)
        (1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): SiLU(inplace=True)
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(240, 10, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(10, 240, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(stochastic_depth): StochasticDepth(p=0.06956521739130435, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
    (2): SqueezeExcitation(
        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)

```

```

        )
    )
    (stochastic_depth): StochasticDepth(p=0.0782608695652174, mode=row)
)
(2): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.08695652173913043, mode=row)
)
(3): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=480, bias=False)
            (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
)
```

```

        (2): SiLU(inplace=True)
    )
(2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
(3): Conv2dNormActivation(
    (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(80, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(stochastic_depth): StochasticDepth(p=0.09565217391304348, mode=row)
)
)
(5): Sequential(
    (0): MBConv(
        (block): Sequential(
            (0): Conv2dNormActivation(
                (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1), bias=False)
                (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
            (1): Conv2dNormActivation(
                (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2), groups=480, bias=False)
                (1): BatchNorm2d(480, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(480, 20, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(20, 480, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
)

```

```

        (stochastic_depth): StochasticDepth(p=0.10434782608695654, mode=row)
    )
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.11304347826086956, mode=row)
)
(2): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
    )
)

```

```

(2): SqueezeExcitation(
    (avgpool): AdaptiveAvgPool2d(output_size=1)
    (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
    (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
    (activation): SiLU(inplace=True)
    (scale_activation): Sigmoid()
)
(3): Conv2dNormActivation(
    (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
(stochastic_depth): StochasticDepth(p=0.12173913043478261, mode=row)
)
(3): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(1, 1), padding=(2,
2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(112, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
(stochastic_depth): StochasticDepth(p=0.13043478260869565, mode=row)
)
)
(6): Sequential(

```

```

(0): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2), padding=(2,
2), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(672, 28, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(28, 672, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(672, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.1391304347826087, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(

```

```

        (avgpool): AdaptiveAvgPool2d(output_size=1)
        (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
        (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
        (activation): SiLU(inplace=True)
        (scale_activation): Sigmoid()
    )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (stochastic_depth): StochasticDepth(p=0.14782608695652175, mode=row)
)
(2): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
            (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): SiLU(inplace=True)
        )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.1565217391304348, mode=row)
)

```

```

(3): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
                (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
                (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1152, 48, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(48, 1152, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
    (3): Conv2dNormActivation(
        (0): Conv2d(1152, 192, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (stochastic_depth): StochasticDepth(p=0.16521739130434784, mode=row)
)
(4): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(192, 1152, kernel_size=(1, 1), stride=(1, 1),
bias=False)
                (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
        (1): Conv2dNormActivation(
            (0): Conv2d(1152, 1152, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=1152, bias=False)
                (1): BatchNorm2d(1152, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
    )
)

```



```

        )
    )
    (stochastic_depth): StochasticDepth(p=0.1826086956521739, mode=row)
)
(1): MBConv(
    (block): Sequential(
        (0): Conv2dNormActivation(
            (0): Conv2d(320, 1920, kernel_size=(1, 1), stride=(1, 1),
bias=False)
                (1): BatchNorm2d(1920, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
        (1): Conv2dNormActivation(
            (0): Conv2d(1920, 1920, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=1920, bias=False)
                (1): BatchNorm2d(1920, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): SiLU(inplace=True)
            )
        (2): SqueezeExcitation(
            (avgpool): AdaptiveAvgPool2d(output_size=1)
            (fc1): Conv2d(1920, 80, kernel_size=(1, 1), stride=(1, 1))
            (fc2): Conv2d(80, 1920, kernel_size=(1, 1), stride=(1, 1))
            (activation): SiLU(inplace=True)
            (scale_activation): Sigmoid()
        )
        (3): Conv2dNormActivation(
            (0): Conv2d(1920, 320, kernel_size=(1, 1), stride=(1, 1),
bias=False)
                (1): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            )
        )
    (stochastic_depth): StochasticDepth(p=0.19130434782608696, mode=row)
)
)
(8): Conv2dNormActivation(
    (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): SiLU(inplace=True)
)
)
(avgpool): AdaptiveAvgPool2d(output_size=1)
(classifier): Sequential(
    (0): Dropout(p=0.2, inplace=True)

```

```
(1): Linear(in_features=1280, out_features=1000, bias=True)
)
)
```

```
[6]: weights.transforms()
```

```
[6]: ImageClassification(
    crop_size=[240]
    resize_size=[255]
    mean=[0.485, 0.456, 0.406]
    std=[0.229, 0.224, 0.225]
    interpolation=InterpolationMode.BILINEAR
)
```

```
[7]: efficientnet_b1_transforms = weights.transforms()
```

```
[8]: train_dataset = datasets.CIFAR10(
    root = "image",
    train = True,
    download = True,
    transform = efficientnet_b1_transforms
)

test_dataset = datasets.CIFAR10(
    root = "image",
    train = False,
    download = True,
    transform = efficientnet_b1_transforms
)
```

Files already downloaded and verified

Files already downloaded and verified

```
[9]: BATCH_SIZE = 32

train_dataloader = DataLoader(train_dataset,batch_size = BATCH_SIZE,shuffle = True)
test_dataloader = DataLoader(test_dataset,batch_size = BATCH_SIZE,shuffle = False)
```

```
[10]: len(train_dataloader)
```

```
[10]: 1563
```

```
[11]: summary(model=model,
    input_size=(32, 3, 128, 128),
    col_names=["input_size","output_size","num_params","trainable"],
    row_settings=["var_names"])
```

Layer (type (var_name))	Output Shape	Param #	Trainable	Input Shape
EfficientNet (EfficientNet)				[32, 3, 128, 128]
[32, 1000]	--		True	[32, 3, 128, 128]
Sequential (features)				[32, 3, 128, 128]
[32, 1280, 4, 4]	--		True	[32, 3, 128, 128]
Conv2dNormActivation (0)				[32, 3, 128, 128]
[32, 32, 64, 64]	--		True	[32, 3, 128, 128]
Conv2d (0)				[32, 3, 128, 128]
[32, 32, 64, 64]	864		True	[32, 32, 64, 64]
BatchNorm2d (1)				[32, 32, 64, 64]
[32, 32, 64, 64]	64		True	[32, 32, 64, 64]
SiLU (2)				[32, 32, 64, 64]
[32, 32, 64, 64]	--		--	[32, 32, 64, 64]
Sequential (1)				[32, 32, 64, 64]
[32, 16, 64, 64]	--		True	[32, 32, 64, 64]
MBConv (0)				[32, 32, 64, 64]
[32, 16, 64, 64]	1,448		True	[32, 16, 64, 64]
MBConv (1)				[32, 16, 64, 64]
[32, 16, 64, 64]	612		True	[32, 16, 64, 64]
Sequential (2)				[32, 16, 64, 64]
[32, 24, 32, 32]	--		True	[32, 16, 64, 64]
MBConv (0)				[32, 16, 64, 64]
[32, 24, 32, 32]	6,004		True	[32, 24, 32, 32]
MBConv (1)				[32, 24, 32, 32]
[32, 24, 32, 32]	10,710		True	[32, 24, 32, 32]
MBConv (2)				[32, 24, 32, 32]
[32, 24, 32, 32]	10,710		True	[32, 24, 32, 32]
Sequential (3)				[32, 24, 32, 32]
[32, 40, 16, 16]	--		True	[32, 24, 32, 32]
MBConv (0)				[32, 24, 32, 32]
[32, 40, 16, 16]	15,350		True	[32, 40, 16, 16]
MBConv (1)				[32, 40, 16, 16]
[32, 40, 16, 16]	31,290		True	[32, 40, 16, 16]
MBConv (2)				[32, 40, 16, 16]
[32, 40, 16, 16]	31,290		True	[32, 40, 16, 16]
Sequential (4)				[32, 40, 16, 16]
[32, 80, 8, 8]	--		True	[32, 40, 16, 16]
MBConv (0)				[32, 40, 16, 16]
[32, 80, 8, 8]	37,130		True	[32, 80, 8, 8]
MBConv (1)				[32, 80, 8, 8]
[32, 80, 8, 8]	102,900		True	[32, 80, 8, 8]
MBConv (2)				[32, 80, 8, 8]

[32, 80, 8, 8]	102,900	True	[32, 80, 8, 8]
MBConv (3)			
[32, 80, 8, 8]	102,900	True	[32, 80, 8, 8]
Sequential (5)			
[32, 112, 8, 8]	--	True	[32, 80, 8, 8]
MBConv (0)			
[32, 112, 8, 8]	126,004	True	[32, 112, 8, 8]
MBConv (1)			
[32, 112, 8, 8]	208,572	True	[32, 112, 8, 8]
MBConv (2)			
[32, 112, 8, 8]	208,572	True	[32, 112, 8, 8]
MBConv (3)			
[32, 112, 8, 8]	208,572	True	[32, 112, 8, 8]
Sequential (6)			
[32, 192, 4, 4]	--	True	[32, 112, 8, 8]
MBConv (0)			
[32, 192, 4, 4]	262,492	True	[32, 192, 4, 4]
MBConv (1)			
[32, 192, 4, 4]	587,952	True	[32, 192, 4, 4]
MBConv (2)			
[32, 192, 4, 4]	587,952	True	[32, 192, 4, 4]
MBConv (3)			
[32, 192, 4, 4]	587,952	True	[32, 192, 4, 4]
MBConv (4)			
[32, 192, 4, 4]	587,952	True	[32, 192, 4, 4]
Sequential (7)			
[32, 320, 4, 4]	--	True	[32, 192, 4, 4]
MBConv (0)			
[32, 320, 4, 4]	717,232	True	[32, 320, 4, 4]
MBConv (1)			
[32, 320, 4, 4]	1,563,600	True	[32, 320, 4, 4]
Conv2dNormActivation (8)			
[32, 1280, 4, 4]	--	True	[32, 320, 4, 4]
Conv2d (0)			
[32, 1280, 4, 4]	409,600	True	[32, 1280, 4, 4]
BatchNorm2d (1)			
[32, 1280, 4, 4]	2,560	True	[32, 1280, 4, 4]
SiLU (2)			
[32, 1280, 4, 4]	--	--	[32, 1280, 4, 4]
AdaptiveAvgPool2d (avgpool)			
[32, 1280, 1, 1]	--	--	[32, 1280]
Sequential (classifier)			
[32, 1000]	--	True	[32, 1280]
Dropout (0)			
[32, 1280]	--	--	[32, 1280]
Linear (1)			
[32, 1000]	1,281,000	True	[32, 1280]

```
=====
=====
Total params: 7,794,184
Trainable params: 7,794,184
Non-trainable params: 0
Total mult-adds (G): 6.01
=====
=====
Input size (MB): 6.29
Forward/backward pass size (MB): 1565.47
Params size (MB): 31.18
Estimated Total Size (MB): 1602.94
=====
```

[12]: # This is transfer learning (feature extraction mode):

```
# Freeze backbone (remain the pretrained features)

# Train only the final classifier to map features → the labels
```

[13]: model.classifier[1] = nn.Linear(in\_features=1280,out\_features=10,bias=True).
 ↪to(device)

[14]: for param in model.features.parameters():
 param.requires\_grad=False
 # print(param)

[15]: from torch.optim.lr\_scheduler import StepLR

```
cost_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(params = model.parameters(), lr = 0.01,
    ↪weight_decay=1e-4)
scheduler = StepLR(optimizer, step_size=5, gamma=0.5)
```

[16]: train\_losses, test\_losses = [], []
train\_accuracies, test\_accuracies = [], []

epochs = 10
for epoch in range(epochs):
 print(f"Epoch: {epoch} \n -----")

 train\_cost, train\_acc = train\_step(train\_dataloader, model, cost\_fn,
 ↪optimizer, accuracy\_fn, device)

 test\_cost, test\_acc = test\_step(test\_dataloader, model, cost\_fn,
 ↪accuracy\_fn, device)

```

scheduler.step()

train_losses.append(train_cost.detach().cpu().item())
test_losses.append(test_cost.detach().cpu().item())
train_accuracies.append(train_acc.detach().cpu().item())
test_accuracies.append(test_acc.detach().cpu().item())

print(f"\nTrain Cost: {train_cost:.4f}, {train_acc:.2f}")

print(f"Test Cost: {test_cost:.4f}, {test_acc:.2f} \n")

```

Epoch: 0

-----

Train Cost: 0.7885, 72.88

Test Cost: 0.6055, 79.57

Epoch: 1

-----

Train Cost: 0.7516, 74.62

Test Cost: 0.5754, 80.43

Epoch: 2

-----

Train Cost: 0.7471, 74.76

Test Cost: 0.5979, 79.51

Epoch: 3

-----

Train Cost: 0.7477, 74.67

Test Cost: 0.5694, 80.54

Epoch: 4

-----

Train Cost: 0.7473, 74.71

Test Cost: 0.5856, 79.88

Epoch: 5

-----

Train Cost: 0.6918, 76.12

Test Cost: 0.5443, 81.19

```
Epoch: 6
-----
Train Cost: 0.6779, 76.67
Test Cost: 0.5303, 81.83
```

```
Epoch: 7
-----
Train Cost: 0.6760, 76.65
Test Cost: 0.5278, 81.66
```

```
Epoch: 8
-----
Train Cost: 0.6716, 76.79
Test Cost: 0.5345, 81.70
```

```
Epoch: 9
-----
Train Cost: 0.6718, 76.54
Test Cost: 0.5324, 81.78
```

```
[17]: # Retraining the last layer is not enough
# Because the model needs to:
# Adapt deeper features
# Ignore irrelevant ImageNet patterns (dog fur, cat eyes, textures, etc.)
```

```
[18]: for param in model.features[-2:].parameters():
    param.requires_grad = True
```

```
[19]: optimizer = torch.optim.Adam(model.parameters(), lr=1e-5)
```

```
[20]: print("\nPhase 2: Fine-tuning the backbone")

epochs = 10
for epoch in range(epochs):
    print(f"Epoch: {epoch} \n -----")

    train_cost, train_acc = train_step(train_dataloader, model, cost_fn,
                                       optimizer, accuracy_fn, device)
```

```

    test_cost, test_acc = test_step(test_dataloader, model, cost_fn, accuracy_fn, device)

scheduler.step()

train_losses.append(train_cost.detach().cpu().item())
test_losses.append(test_cost.detach().cpu().item())
train_accuracies.append(train_acc.detach().cpu().item())
test_accuracies.append(test_acc.detach().cpu().item())

print(f"\nTrain Cost: {train_cost:.4f}, {train_acc:.2f}")

print(f"Test Cost: {test_cost:.4f}, {test_acc:.2f} \n")

```

Phase 2: Fine-tuning the backbone

Epoch: 0

-----

Train Cost: 0.6126, 78.72

Test Cost: 0.4811, 83.61

Epoch: 1

-----

Train Cost: 0.5616, 80.41

Test Cost: 0.4561, 84.44

Epoch: 2

-----

Train Cost: 0.5321, 81.31

Test Cost: 0.4534, 84.72

Epoch: 3

-----

Train Cost: 0.5153, 81.93

Test Cost: 0.4266, 85.50

Epoch: 4

-----

Train Cost: 0.4986, 82.40

Test Cost: 0.4162, 85.88

Epoch: 5

-----

```
Train Cost: 0.4817, 83.04
Test Cost: 0.4105, 86.09
```

```
Epoch: 6
```

```
-----
```

```
Train Cost: 0.4637, 83.86
Test Cost: 0.4019, 86.45
```

```
Epoch: 7
```

```
-----
```

```
Train Cost: 0.4521, 84.11
Test Cost: 0.3990, 86.45
```

```
Epoch: 8
```

```
-----
```

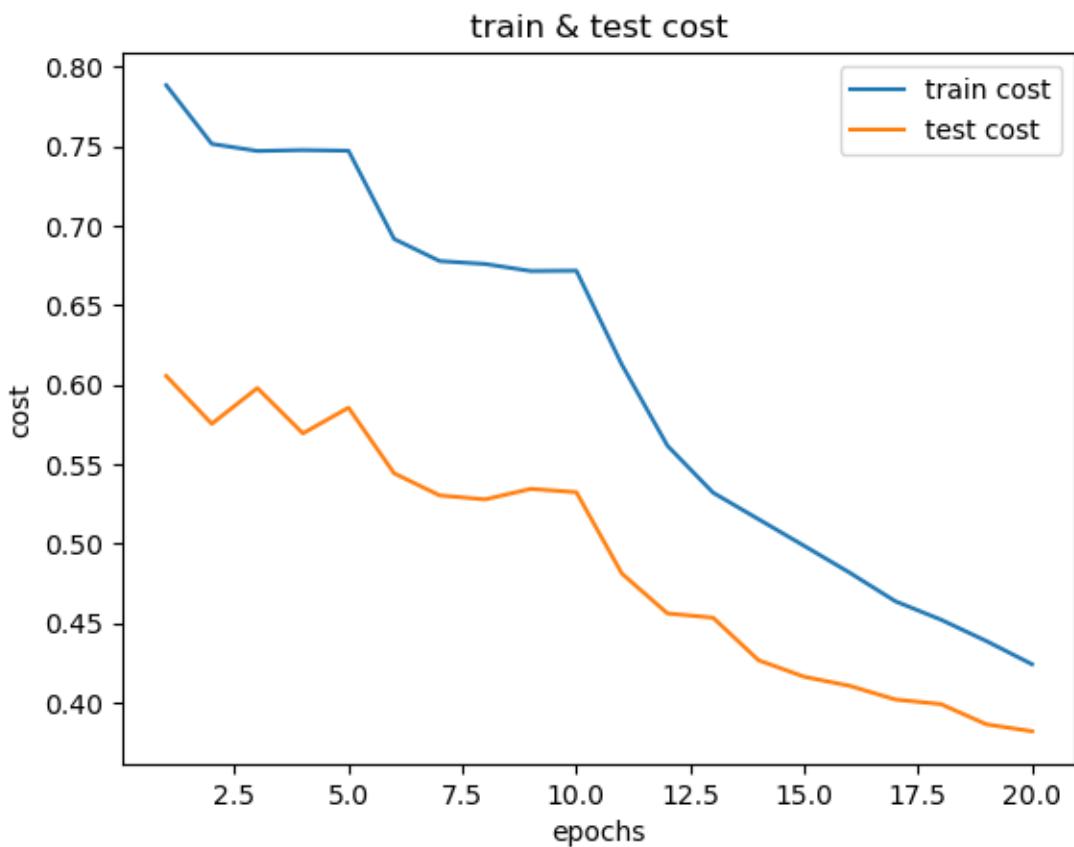
```
Train Cost: 0.4387, 84.58
Test Cost: 0.3862, 87.01
```

```
Epoch: 9
```

```
-----
```

```
Train Cost: 0.4241, 85.10
Test Cost: 0.3819, 87.00
```

```
[21]: plt.plot(range(1,len(train_losses) + 1), train_losses, label = "train cost")
plt.plot(range(1,len(test_losses) + 1), test_losses, label = "test cost")
plt.title("train & test cost")
plt.xlabel("epochs")
plt.ylabel("cost")
plt.legend()
plt.show()
```



## 0.2 save model

```
[25]: # torch.save(model.state_dict(), "efficientnet_cifar10.pth")
```